

# CoCoA<sup>+</sup>: Adding vs. Averaging in Distributed Optimization

Martin Takáč



joint work with

- Chenxin Ma
- Virginia Smith
- Martin Jaggi
- Michael I. Jordan
- Peter Richtrik

- Problem Formulation
- Serial and Parallel Coordinate Descent Method (CDM)
- Distributed CDM
- Original CoCoA Framework
- CoCoA<sup>+</sup> Framework
- Computation vs. Communication Trade-off
- Spark
- Some Numerical Experiments

# The Problem - Regularized Empirical Loss Minimization

Let  $\{(x_i, y_i)\}_{i=1}^n$  be our training data,  $x_i \in \mathbf{R}^d$  and  $y_i \in \mathbf{R}$ .

$$\min_{w \in \mathbf{R}^d} \left[ P(w) := \frac{1}{n} \sum_{i=1}^n \ell_i(w^T x_i) + \frac{\lambda}{2} \|w\|^2 \right] \quad (\text{P})$$

where

- $\lambda > 0$  is a regularization parameter
- $\ell_i(\cdot)$  is convex loss function which can depend on the label  $y_i$

*Examples:*

- Logistic loss:  $\ell_i(\zeta) = \log(1 + \exp(-y_i \zeta))$
- Hinge loss:  $\ell_i(\zeta) = \max\{0, 1 - y_i \zeta\}$

# The Problem - Regularized Empirical Loss Minimization

Let  $\{(x_i, y_i)\}_{i=1}^n$  be our training data,  $x_i \in \mathbf{R}^d$  and  $y_i \in \mathbf{R}$ .

$$\min_{w \in \mathbf{R}^d} \left[ P(w) := \frac{1}{n} \sum_{i=1}^n \ell_i(w^T x_i) + \frac{\lambda}{2} \|w\|^2 \right] \quad (\text{P})$$

where

- $\lambda > 0$  is a regularization parameter
- $\ell_i(\cdot)$  is convex loss function which can depend on the label  $y_i$

*Examples:*

- Logistic loss:  $\ell_i(\zeta) = \log(1 + \exp(-y_i \zeta))$
- Hinge loss:  $\ell_i(\zeta) = \max\{0, 1 - y_i \zeta\}$

## The dual problem

$$\max_{\alpha \in \mathbf{R}^n} \left[ D(\alpha) := -\frac{\lambda}{2} \|A\alpha\|^2 - \frac{1}{n} \sum_{i=1}^n \ell_i^*(-\alpha_i) \right] \quad (\text{D})$$

where  $A = \frac{1}{\lambda n} X^T$  and  $X^T = [x_1, x_2, \dots, x_n] \in \mathbf{R}^{d \times n}$

- $\ell_i^*$  is convex conjugate of  $\ell_i$
- wlog  $\|x_i\| \leq 1$

## Primal-Dual mapping

For any  $\alpha \in \text{dom}(D)$  we can define

$$w(\alpha) := A\alpha \quad (1)$$

From strong duality we have that  $w^* = w(\alpha^*)$  is optimal to (P) if  $\alpha^*$  is optimal solution to (D).

## Gap function

$$G(\alpha) = P(w(\alpha)) - D(\alpha)$$

# The Setting & Challenges

- The size of matrix  $A$  is huge (e.g. TBs of data)
- We want to use many nodes of computer cluster (or cloud) to speed-up the computation

## Challenges

- **distributed data:** no single machine can load the whole instance
- **expensive communication:**

	latency
RAM	100 nanoseconds
standard network connection	250,000 nanoseconds

- **unreliable nodes:** we assume that the node can die at any point during the computation (we want to have fault tolerant solution)

# The Serial/Parallel/Distributed SDCA Algorithm

## Serial **S**tochastic **D**ual **C**oordinate **A**scent

choose  $\alpha^{(0)} \in \mathbf{R}^n$

repeat

$$\alpha^{(t+1)} = \alpha^{(t)}$$

pick a random coordinate  $i \in \{1, \dots, n\}$

compute the update:  $h_t^i(\alpha^{(t)}) := \arg \max_h D(\alpha^{(t)} + h e_i)$

apply the update:  $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + h_t^i(\alpha^{(t)}) e_i$

# The Serial/Parallel/Distributed SDCA Algorithm

## Parallel Stochastic Dual Coordinate Ascent

choose  $\alpha^{(0)} \in \mathbf{R}^n$

repeat

$$\alpha^{(t+1)} = \alpha^{(t)}$$

pick a random coordinate  $i \in \{1, \dots, n\}$

pick a random subset  $S \subset \{1, \dots, n\}$  with  $|S| = H$

for each  $i \in S$  in **parallel** do

compute the update:  $h_t^i(\alpha^{(t)}) := \arg \max_h D(\alpha^{(t)} + he_i)$

apply the update:  $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + h_t^i(\alpha^{(t)})e_i$

apply the update:  $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + \frac{1}{H} \sum_{i \in S} h_t^i(\alpha^{(t)})e_i$



# The Serial/Parallel/Distributed SDCA Algorithm

- assume we have  $K$  nodes (computers) each with parallel processing power
- we **partition** the coordinates  $\{1, 2, \dots, n\}$  into  $K$  **balanced** sets  $\mathcal{P}_1, \dots, \mathcal{P}_K$   
 $\forall k \in \{1, \dots, K\}$  we have  $|\mathcal{P}_k| = \frac{n}{K}$

## Distributed Stochastic Dual Coordinate Ascent

choose  $\alpha^{(0)} \in \mathbf{R}^n$

repeat

$$\alpha^{(t+1)} = \alpha^{(t)}$$

for each **computer**  $k \in \{1, \dots, K\}$  in **parallel** do

pick a random subset  $S \subset \{1, \dots, n\}$  with  $|S| = H$

pick a random subset  $S_k \subset \mathcal{P}_k$  with  $|S_k| = H \leq \frac{n}{K}$

for each  $i \in S_k$  in **parallel** do

compute the update:  $h_t^i(\alpha^{(t)}) := \arg \max_h D(\alpha^{(t)} + h e_i)$


apply the update:  $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + \frac{1}{H} \sum_{i \in S} h_t^i(\alpha^{(t)}) e_i$

apply the update:  $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + \frac{1}{KH} \sum_{k \in \{1, \dots, K\}} \sum_{i \in S_k} h_t^i(\alpha^{(t)}) e_i$

The distributed algorithm can need (in the worst case) the **same number of iterations as a serial one!**


# Distributed CDM

Illustration:  $K = 4$  and  $H = 2$




# Distributed CDM

Illustration:  $K = 4$  and  $H = 2$




# Distributed CDM

Illustration:  $K = 4$  and  $H = 2$



# Distributed CDM

Illustration:  $K = 4$  and  $H = 2$




# Disadvantages of Distributed CDM

- we cannot choose  $H > |\mathcal{P}_k|!$
- the computation of step is very easy (usually close form or a bit complicated 1D problem)
- after taking  $H$  steps, usually the objective function doesn't change much
- it is almost **impossible** to balance computation and communication

# Data Distribution

Vector  $\alpha$  and columns of matrix  $A$  are partitioned according to  $\{\mathcal{P}_k\}_{k=1}^K$ .



**Notation:** For  $k \in \{1, 2, \dots, K\}$  we use  $\alpha_k \in \mathbf{R}^{|\mathcal{P}_k|}$  is a subvector of  $\alpha$ . Vector  $\alpha_{[k]} \in \mathbf{R}^n$  is a vector obtained from vector  $\alpha$  by setting all coordinates  $\notin \mathcal{P}_k$  to zero.

**Example:**  $\alpha_1 = (*, *, *, *)^T$ ,  $\alpha_{[1]} = (*, *, *, *, 0, 0, \dots, 0)^T$ .

# Local Problem

## CoCoA subproblem

At iteration  $t$  at node  $k$

$$\begin{aligned}(\Delta\alpha^*)_{[k]}^{(t)} &= \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} D(\alpha^{(t)} + \Delta\alpha_{[k]}) \\ &= \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \left( -\frac{\lambda}{2} \|A(\alpha^{(t)} + \Delta\alpha_{[k]})\|^2 - \frac{1}{n} \sum_{i=1}^n \ell_i^*(-(\alpha^{(t)} + \Delta\alpha_{[k]})_i) \right)\end{aligned}$$

- we cannot solve the subproblem as it depends on  $\alpha^{(t)}$  and  $A$



# Local Problem

## CoCoA subproblem

At iteration  $t$  at node  $k$

$$\begin{aligned}(\Delta\alpha^*)_{[k]}^{(t)} &= \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} D(\alpha^{(t)} + \Delta\alpha_{[k]}) \\ &= \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \left( -\frac{\lambda}{2} \|A(\alpha^{(t)} + \Delta\alpha_{[k]})\|^2 - \frac{1}{n} \sum_{i=1}^n \ell_i^*(-(\alpha^{(t)} + \Delta\alpha_{[k]})_i) \right)\end{aligned}$$

- we cannot solve the subproblem as it depends on  $\alpha^{(t)}$  and  $A$
- if we know  $w^{(t)} = A\alpha^{(t)}$  then

$$(\Delta\alpha^*)_{[k]}^{(t)} = \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \left( -\frac{\lambda}{2} \|w^{(t)} + A\Delta\alpha_{[k]}\|^2 - \frac{1}{n} \sum_{i \in \mathcal{P}_k} \ell_i^*(-(\alpha^{(t)} + \Delta\alpha_{[k]})_i) \right)$$

- if we know  $w^{(t)}$  we can compute  $(\Delta\alpha^*)_{[k]}^{(t)}$

# The CoCoA Framework

## Communication-Efficient Distributed Dual Coordinate Ascent

**Input:**  $T \geq 1$

**Data:**  $\{(x_i, y_i)\}_{i=1}^n$  distributed over  $K$  machines

**Initialize:**  $\alpha_{[k]}^{(0)} \leftarrow 0$  for all machines  $k$ , and  $w^{(0)} \leftarrow 0$

**for**  $t = 1, 2, \dots, T$

**for all machines**  $k = 1, 2, \dots, K$  **in parallel**

Solve local problem approximately to obtain  $\Delta\alpha_{[k]}$

computation

$$\alpha_{[k]}^{(t)} \leftarrow \alpha_{[k]}^{(t-1)} + \frac{1}{K} \Delta\alpha_{[k]}$$

$$\Delta w_k \leftarrow \frac{1}{K} A \Delta\alpha_{[k]}$$

$$\text{reduce } w^{(t)} \leftarrow w^{(t-1)} + \sum_{k=1}^K \Delta w_k$$

communication

- The performance of this methods (in worst case) can be the same as if we **randomly pick**  $k$  and solve corresponding subproblem and replace  $\frac{1}{K}$  by 1
- How **accurately** do we need to solve the local sub-problem?
- How to change the local problem to avoid averaging (e.g. just to add local solutions)?
- Can we prove it will be better?

# Smarter Subproblem

## Local Subproblem for CoCoA<sup>+</sup>

$$\max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}; w^{(t)}) \quad (2)$$

where

$$\begin{aligned} \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}; w^{(t)}) = & -\frac{1}{n} \sum_{i \in \mathcal{P}_k} \ell_i^*(-(\alpha_{[k]}^{(t)} + \Delta\alpha_{[k]})_i) \\ & - \frac{1}{K} \frac{\lambda}{2} \|w^{(t)}\|^2 - \lambda (w^{(t)})^T A \Delta\alpha_{[k]} \\ & - \frac{\lambda}{2} \sigma' \|A \Delta\alpha_{[k]}\|^2. \end{aligned}$$

Compare with:

$$\max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \left( -\frac{\lambda}{2} \|w^{(t)} + A \Delta\alpha_{[k]}\|^2 - \frac{1}{n} \sum_{i \in \mathcal{P}_k} \ell_i^*(-(\alpha^{(t)} + \Delta\alpha_{[k]})_i) \right) \quad (3)$$

If  $\sigma' = 1$  then the optimal solutions of (2) and (3) coincide.

# The CoCoA<sup>+</sup> Framework

## Communication-Efficient Distributed Dual Coordinate Ascent

**Input:**  $T \geq 1$ ,  $\gamma \in [\frac{1}{K}, 1]$ ,  $\sigma' \in [1, \infty)$

**Data:**  $\{(x_i, y_i)\}_{i=1}^n$  distributed over  $K$  machines

**Initialize:**  $\alpha_{[k]}^{(0)} \leftarrow 0$  for all machines  $k$ , and  $w^{(0)} \leftarrow 0$

**for**  $t = 1, 2, \dots, T$

**for all machines**  $k = 1, 2, \dots, K$  **in parallel**

        approximately  $\max \mathcal{G}^{\sigma'}(\Delta\alpha_{[k]}; w^{(t)})$  to obtain  $\Delta\alpha_{[k]}$  computation

$$\alpha_{[k]}^{(t)} \leftarrow \alpha_{[k]}^{(t-1)} + \gamma \Delta\alpha_{[k]}$$

$$\Delta w_k \leftarrow \gamma A \Delta\alpha_{[k]}$$

    reduce  $w^{(t)} \leftarrow w^{(t-1)} + \sum_{k=1}^K \Delta w_k$  communication

- If  $\gamma = \frac{1}{K}$  we obtain CoCoA
- If  $\gamma = \frac{1}{K}$  then  $\sigma' = 1$  is "safe" value
- What about another values of  $\gamma$ ? (we want  $\gamma = 1$ )

# CoCoA<sup>+</sup> Parameters - $\sigma'$ and $\gamma$

- $\sigma'$  measures the difficulty of the given data partition
- it must be chosen not smaller than

$$\sigma' \geq \sigma'_{\min} \stackrel{\text{def}}{=} \gamma \max_{\alpha \in \mathbf{R}^n} \frac{\|A\alpha\|^2}{\sum_{k=1}^K \|A\alpha_{[k]}\|^2} \quad (4)$$

## Lemma

For any  $\alpha \in \mathbf{R}^n$  ( $\alpha \neq \mathbf{0}$ ) we have

$$\frac{\|A\alpha\|^2}{\sum_{k=1}^K \|A\alpha_{[k]}\|^2} \leq K$$

- We can take the safe value  $\sigma' = K \cdot \gamma$   
Again: if  $\gamma = \frac{1}{K}$  then  $\sigma' = K \cdot \frac{1}{K} = 1$  is a safe value
- New: if  $\gamma = 1$  then  $\sigma' = K \cdot 1 = K$  is a safe value

# How Accurately?

## Assumption: $\Theta$ -approximate solution

We assume that there exists  $\Theta \in [0, 1]$  such that  $\forall k \in [K]$ , the local solver at any iteration  $t$  produces a **(possibly) randomized** approximate solution  $\Delta\alpha_{[k]}$ , which satisfies

$$\mathbf{E}[\mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}^*, w) - \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}, w)] \leq \Theta \left( \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}^*, w) - \mathcal{G}_k^{\sigma'}(\mathbf{0}, w) \right), \quad (5)$$

where

$$\Delta\alpha^* \in \arg \min_{\Delta\alpha \in \mathbf{R}^n} \sum_{k=1}^K \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}, w). \quad (6)$$

- because the subproblem is **not really** what one wants to solve, therefore in practise  $\Theta \approx 0.9$  (depending on the cluster and problem)
- what about convergence guarantees?
- how to get  $\Theta$  approximate solution?

# Iteration Complexity - Smooth Loss

## Theorem

Assume the loss functions  $\ell_i$  are  $(1/\mu)$ -smooth, for  $i \in \{1, 2, \dots, n\}$ . We define

$$\sigma_k \stackrel{\text{def}}{=} \max_{\alpha_{[k]} \in \mathbf{R}^n} \frac{\|A\alpha_{[k]}\|^2}{\|\alpha_{[k]}\|^2} \leq |\mathcal{P}_k| \quad (7)$$

and  $\sigma_{\max} = \max_{k \in [K]} \sigma_k$ .

Then after  $T$  iterations of CoCoA<sup>+</sup>, with

$$T \geq \frac{1}{\gamma(1-\Theta)} \frac{\lambda\mu n + \sigma_{\max}\sigma'}{\lambda\mu n} \log \frac{1}{\epsilon},$$

it holds that  $\mathbf{E}[D(\alpha^*) - D(\alpha^T)] \leq \epsilon$ .

Furthermore, after  $T$  iterations with

$$T \geq \frac{1}{\gamma(1-\Theta)} \frac{\lambda\mu n + \sigma_{\max}\sigma'}{\lambda\mu n} \log \left( \frac{1}{\gamma(1-\Theta)} \frac{\lambda\mu n + \sigma_{\max}\sigma'}{\lambda\mu n} \frac{1}{\epsilon} \right),$$

we have the expected duality gap

$$\mathbf{E}[P(w(\alpha^{(T)})) - D(\alpha^{(T)})] \leq \epsilon.$$

# Averaging vs. Adding

The leading term is  $\frac{1}{\gamma(1-\Theta)} \frac{\lambda\mu n + \sigma_{\max}\sigma'}{\lambda\mu n}$ . Let us assume that  $\forall k : |\mathcal{P}_k| = \frac{n}{K}$

## Averaging

$$\gamma = \frac{1}{K}$$
$$\sigma' = 1$$

$$\frac{\textcolor{red}{K}}{1-\Theta} \frac{\lambda\mu n + \frac{n}{K}}{\lambda\mu n}$$

$$\frac{1}{1-\Theta} \frac{\lambda\mu \textcolor{red}{K} + 1}{\lambda\mu}$$

## Adding

$$\gamma = 1$$
$$\sigma' = K$$

$$\frac{1}{1-\Theta} \frac{\lambda\mu n + \frac{n}{K} \textcolor{red}{K}}{\lambda\mu n}$$

$$\frac{1}{1-\Theta} \frac{\lambda\mu + 1}{\lambda\mu}$$

Note: this is in the worst case (for the worst case example)



# Iteration Complexity - General Convex Loss

## Theorem

Consider CoCoA<sup>+</sup> starting with  $\alpha^0 = \mathbf{0} \in \mathbf{R}^n$  and  $\forall i \in \{1, 2, \dots, n\} : \ell_i(\cdot)$  be  $L$ -Lipschitz continuous and  $\epsilon > 0$  be the desired duality gap. Then after  $T$  iterations, where

$$\begin{aligned} T &\geq T_0 + \max\left\{\left\lceil \frac{1}{\gamma(1-\Theta)} \right\rceil, \frac{4L^2\sigma\sigma'}{\lambda n^2\epsilon\gamma(1-\Theta)}\right\}, \\ T_0 &\geq t_0 + \left(\frac{2}{\gamma(1-\Theta)} \left(\frac{8L^2\sigma\sigma'}{\lambda n^2\epsilon} - 1\right)\right)_+, \\ t_0 &\geq \max(0, \left\lceil \frac{1}{\gamma(1-\Theta)} \log\left(\frac{2\lambda n^2(D(\alpha^*) - D(\alpha^0))}{4L^2\sigma\sigma'}\right) \right\rceil), \end{aligned}$$

we have that the expected duality gap satisfies  $\mathbf{E}[P(w(\bar{\alpha})) - D(\bar{\alpha})] \leq \epsilon$ , at the averaged iterate

$$\bar{\alpha} := \frac{1}{T-T_0} \sum_{t=T_0+1}^{T-1} \alpha^{(t)},$$

where  $\sigma = \sum_{k=1}^K |\mathcal{P}_k| \sigma_k$ .

# SDCA as a Local Solver

## SDCA

- 1: **Input:**  $\alpha_{[k]}, w = w(\alpha)$
- 2: **Data:** Local  $\{(x_i, y_i)\}_{i \in \mathcal{P}_k}$
- 3: **Initialize:**  $\Delta\alpha_{[k]}^0 = 0 \in \mathbb{R}^n$
- 4: **for**  $h = 0, 1, \dots, H - 1$  **do**
- 5:   choose  $i \in \mathcal{P}_k$  uniformly at random
- 6:    $\delta_i^* = \arg \max_{\delta_i \in \mathbb{R}} \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}^h + \delta_i e_i, w)$
- 7:    $\Delta\alpha_{[k]}^{(h+1)} = \Delta\alpha_{[k]}^{(h)} + \delta_i^* e_i$
- 8: **end for**
- 9: **Output:**  $\Delta\alpha_{[k]}^{(H)}$

## Theorem

Assume the functions  $\ell_i$  are  $(1/\mu)$ -smooth for  $i \in \{1, 2, \dots, n\}$ . If

$$H \geq n_k \frac{\sigma' + \lambda n \mu}{\lambda n \mu} \log \frac{1}{\Theta} \quad (8)$$

then SDCA will produce a  $\Theta$ -approximate solution.

# Total Runtime

- To get  $\epsilon$  accuracy we need

$$\mathcal{O}\left(\frac{1}{1 - \Theta} \log \frac{1}{\epsilon}\right)$$

- Recall  $\Theta = \left(1 - \frac{\lambda n \gamma}{1 + \lambda n \gamma} \frac{K}{n}\right)^H$





Let

- $\tau_o$  be the duration of communication per iteration
- $\tau_c$  be the duration of **ONE** coordinate update during the inner iteration

## Total runtime

$$\mathcal{O}\left(\frac{1}{1 - \Theta} (\tau_o + H \tau_c)\right) = \mathcal{O}\left(\frac{1}{1 - \Theta} \left(1 + \underbrace{H \frac{\tau_c}{\tau_o}}_{r_{c/o}}\right)\right)$$

$$H(\tau_c/\tau_o), \Theta(\tau_c/\tau_o)$$



## Apache Spark


- is a fast and general engine for large-scale data processing
- runs on Hadoop, Mesos, standalone, or in the cloud. It can access diverse data sources including HDFS, Cassandra, HBase, S3
- **is slower than our C++ code** for CoCoA<sup>+</sup>
- we run it on Amazon Elastic Compute Cloud (Amazon EC2)

# Numerical Experiments

## Datasets

Dataset	Training ( $n$ )	Features ( $d$ )	$nnz/(n \cdot d)$	$\lambda$	Workers ( $K$ )
cov	522,911	54	22.22%	$1e-6$	4
rcv1	677,399	47,236	0.16%	$1e-6$	8
imagenet	32,751	160,000	100%	$1e-5$	32

# Dependence of Primal Suboptimality on $H$



# Comparison with Different Algorithms

- COCOA



# Comparison with Different Algorithms

- COCOA

- minibatch-CD

Tianbao Yang. *Trading Computation for Communication: Distributed Stochastic Dual Coordinate Ascent*. In NIPS 2013.

Martin Takáč, Avleen Bijral, Peter Richtárik, and Nathan Srebro. *Mini-Batch Primal and Dual Methods for SVMs*. In ICML, March 2013.

# Comparison with Different Algorithms

- COCOA
- minibatch-CD  
Tianbao Yang. *Trading Computation for Communication: Distributed Stochastic Dual Coordinate Ascent*. In NIPS 2013.  
Martin Takáč, Avleen Bijral, Peter Richtárik, and Nathan Srebro. *Mini-Batch Primal and Dual Methods for SVMs*. In ICML, March 2013.
- local-SGD  
Shai Shalev-Shwartz, Yoram Singer, Nathan Srebro, and Andrew Cotter. *Pegasos: Primal Estimated Sub-Gradient Solver for SVM*. Mathematical Programming, 127(1):3–30, October 2010.

# Comparison with Different Algorithms

- COCOA

- minibatch-CD

Tianbao Yang. *Trading Computation for Communication: Distributed Stochastic Dual Coordinate Ascent*. In NIPS 2013.




Martin Takáč, Avleen Bijral, Peter Richtárik, and Nathan Srebro. *Mini-Batch Primal and Dual Methods for SVMs*. In ICML, March 2013.




- local-SGD

Shai Shalev-Shwartz, Yoram Singer, Nathan Srebro, and Andrew Cotter. *Pegasos: Primal Estimated Sub-Gradient Solver for SVM*. Mathematical Programming, 127(1):3–30, October 2010.





- batch-SGD

Shai Shalev-Shwartz, Yoram Singer, Nathan Srebro, and Andrew Cotter. *Pegasos: Primal Estimated Sub-Gradient Solver for SVM*. Mathematical Programming, 127(1):3–30, October 2010.












# CoCoA vs. CoCoA<sup>+</sup>



# CoCoA vs. CoCoA<sup>+</sup>





# Scaling up





# Effect of $\sigma'$



# References

- ① Chenxin Ma, Virginia Smith, Martin Jaggi, Michael I. Jordan, Peter Richtárik and Martin Takáč: *Adding vs. Averaging in Distributed Primal-Dual Optimization*, ICML 2015.
- ② Martin Jaggi, Virginia Smith, Martin Takáč, Jonathan Terhorst, Thomas Hofmann and Michael I. Jordan: *Communication-Efficient Distributed Dual Coordinate Ascent*, NIPS 2014.
- ③ Richtárik, P. and Takáč, M.: *On optimal probabilities in stochastic coordinate descent methods*, arXiv:1310.3438, 2013.
- ④ Richtárik, P. and Takáč, M.: *Parallel coordinate descent methods for big data optimization*, arXiv:1212.0873, 2012.
- ⑤ Richtárik, P. and Takáč, M.: *Iteration complexity of randomized block-coordinate descent methods for minimizing a composite function*, Mathematical Programming, 2012.
- ⑥ Takáč, M., Bijral, A., Richtárik, P. and Srebro, N.: *Mini-batch primal and dual methods for SVMs*, In ICML, 2013.
- ⑦ Qu, Z., Richtárik, P. and Zhang, T.: *Randomized dual coordinate ascent with arbitrary sampling*, arXiv:1411.5873, 2014.
- ⑧ Qu, Z., Richtárik, P., Takáč, M. and Fercoq, O.: *SDNA: Stochastic Dual Newton Ascent for Empirical Risk Minimization*, arXiv:1502.02268, 2015.
- ⑨ Tappenden, R., Takáč, M. and Richtárik, P.: *On the Complexity of Parallel Coordinate Descent*, arXiv: 1503.03033, 2015.