

CoCoA⁺: Adding vs. Averaging in Distributed Optimization

Martin Takáč



joint work with

- Chenxin Ma
- Virginia Smith
- Martin Jaggi
- Michael I. Jordan
- Peter Richtrik

- Problem Formulation
- Serial and Parallel Coordinate Descent Method (CDM)
- Distributed CDM
- Original CoCoA Framework
- CoCoA⁺ Framework
- Computation vs. Communication Trade-off
- Spark
- Some Numerical Experiments

The Problem - Regularized Empirical Loss Minimization

Let $\{(x_i, y_i)\}_{i=1}^n$ be our training data, $x_i \in \mathbf{R}^d$ and $y_i \in \mathbf{R}$.

$$\min_{w \in \mathbf{R}^d} \left[P(w) := \frac{1}{n} \sum_{i=1}^n \ell_i(w^T x_i) + \frac{\lambda}{2} \|w\|^2 \right] \quad (\text{P})$$

where

- $\lambda > 0$ is a regularization parameter
- $\ell_i(\cdot)$ is convex loss function which can depend on the label y_i

Examples:

- Logistic loss: $\ell_i(\zeta) = \log(1 + \exp(-y_i \zeta))$
- Hinge loss: $\ell_i(\zeta) = \max\{0, 1 - y_i \zeta\}$

The Problem - Regularized Empirical Loss Minimization

Let $\{(x_i, y_i)\}_{i=1}^n$ be our training data, $x_i \in \mathbf{R}^d$ and $y_i \in \mathbf{R}$.

$$\min_{w \in \mathbf{R}^d} \left[P(w) := \frac{1}{n} \sum_{i=1}^n \ell_i(w^T x_i) + \frac{\lambda}{2} \|w\|^2 \right] \quad (\text{P})$$

where

- $\lambda > 0$ is a regularization parameter
- $\ell_i(\cdot)$ is convex loss function which can depend on the label y_i

Examples:

- Logistic loss: $\ell_i(\zeta) = \log(1 + \exp(-y_i \zeta))$
- Hinge loss: $\ell_i(\zeta) = \max\{0, 1 - y_i \zeta\}$

The dual problem

$$\max_{\alpha \in \mathbf{R}^n} \left[D(\alpha) := -\frac{\lambda}{2} \|A\alpha\|^2 - \frac{1}{n} \sum_{i=1}^n \ell_i^*(-\alpha_i) \right] \quad (\text{D})$$

where $A = \frac{1}{\lambda n} X^T$ and $X^T = [x_1, x_2, \dots, x_n] \in \mathbf{R}^{d \times n}$

- ℓ_i^* is convex conjugate of ℓ_i
- wlog $\|x_i\| \leq 1$

Primal-Dual mapping

For any $\alpha \in \text{dom}(D)$ we can define

$$w(\alpha) := A\alpha \quad (1)$$

From strong duality we have that $w^* = w(\alpha^*)$ is optimal to (P) if α^* is optimal solution to (D).

Gap function

$$G(\alpha) = P(w(\alpha)) - D(\alpha)$$

The Setting & Challenges

- The size of matrix A is huge (e.g. TBs of data)
- We want to use many nodes of computer cluster (or cloud) to speed-up the computation

Challenges

- **distributed data:** no single machine can load the whole instance
- **expensive communication:**

	latency
RAM	100 nanoseconds
standard network connection	250,000 nanoseconds

- **unreliable nodes:** we assume that the node can die at any point during the computation (we want to have fault tolerant solution)

The Serial/Parallel/Distributed SDCA Algorithm

Serial **S**tochastic **D**ual **C**oordinate **A**scent

choose $\alpha^{(0)} \in \mathbf{R}^n$

repeat

$$\alpha^{(t+1)} = \alpha^{(t)}$$

pick a random coordinate $i \in \{1, \dots, n\}$

compute the update: $h_t^i(\alpha^{(t)}) := \arg \max_h D(\alpha^{(t)} + h e_i)$

apply the update: $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + h_t^i(\alpha^{(t)}) e_i$

The Serial/Parallel/Distributed SDCA Algorithm

Parallel Stochastic Dual Coordinate Ascent

choose $\alpha^{(0)} \in \mathbf{R}^n$

repeat

$$\alpha^{(t+1)} = \alpha^{(t)}$$

pick a random coordinate $i \in \{1, \dots, n\}$

pick a random subset $S \subset \{1, \dots, n\}$ with $|S| = H$

for each $i \in S$ in **parallel** do

compute the update: $h_t^i(\alpha^{(t)}) := \arg \max_h D(\alpha^{(t)} + he_i)$

apply the update: $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + h_t^i(\alpha^{(t)})e_i$

apply the update: $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + \frac{1}{H} \sum_{i \in S} h_t^i(\alpha^{(t)})e_i$

The Serial/Parallel/Distributed SDCA Algorithm

- assume we have K nodes (computers) each with parallel processing power
- we **partition** the coordinates $\{1, 2, \dots, n\}$ into K **balanced** sets $\mathcal{P}_1, \dots, \mathcal{P}_K$
 $\forall k \in \{1, \dots, K\}$ we have $|\mathcal{P}_k| = \frac{n}{K}$

Distributed Stochastic Dual Coordinate Ascent

choose $\alpha^{(0)} \in \mathbf{R}^n$

repeat

$$\alpha^{(t+1)} = \alpha^{(t)}$$

for each **computer** $k \in \{1, \dots, K\}$ in **parallel** do

pick a random subset $S \subset \{1, \dots, n\}$ with $|S| = H$

pick a random subset $S_k \subset \mathcal{P}_k$ with $|S_k| = H \leq \frac{n}{K}$

for each $i \in S_k$ in **parallel** do

compute the update: $h_t^i(\alpha^{(t)}) := \arg \max_h D(\alpha^{(t)} + h e_i)$


apply the update: $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + \frac{1}{H} \sum_{i \in S} h_t^i(\alpha^{(t)}) e_i$

apply the update: $\alpha_i^{(t+1)} = \alpha_i^{(t+1)} + \frac{1}{KH} \sum_{k \in \{1, \dots, K\}} \sum_{i \in S_k} h_t^i(\alpha^{(t)}) e_i$

The distributed algorithm can need (in the worst case) the **same number of iterations as a serial one!**


Distributed CDM

Illustration: $K = 4$ and $H = 2$




Distributed CDM

Illustration: $K = 4$ and $H = 2$




Distributed CDM

Illustration: $K = 4$ and $H = 2$



Distributed CDM

Illustration: $K = 4$ and $H = 2$




Disadvantages of Distributed CDM

- we cannot choose $H > |\mathcal{P}_k|!$
- the computation of step is very easy (usually close form or a bit complicated 1D problem)
- after taking H steps, usually the objective function doesn't change much
- it is almost **impossible** to balance computation and communication

Data Distribution

Vector α and columns of matrix A are partitioned according to $\{\mathcal{P}_k\}_{k=1}^K$.



Notation: For $k \in \{1, 2, \dots, K\}$ we use $\alpha_k \in \mathbf{R}^{|\mathcal{P}_k|}$ is a subvector of α . Vector $\alpha_{[k]} \in \mathbf{R}^n$ is a vector obtained from vector α by setting all coordinates $\notin \mathcal{P}_k$ to zero.

Example: $\alpha_1 = (*, *, *, *)^T$, $\alpha_{[1]} = (*, *, *, *, 0, 0, \dots, 0)^T$.

Local Problem

CoCoA subproblem

At iteration t at node k

$$\begin{aligned}(\Delta\alpha^*)_{[k]}^{(t)} &= \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} D(\alpha^{(t)} + \Delta\alpha_{[k]}) \\ &= \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \left(-\frac{\lambda}{2} \|A(\alpha^{(t)} + \Delta\alpha_{[k]})\|^2 - \frac{1}{n} \sum_{i=1}^n \ell_i^*(-(\alpha^{(t)} + \Delta\alpha_{[k]})_i) \right)\end{aligned}$$

- we cannot solve the subproblem as it depends on $\alpha^{(t)}$ and A

Local Problem

CoCoA subproblem

At iteration t at node k

$$\begin{aligned}(\Delta\alpha^*)_{[k]}^{(t)} &= \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} D(\alpha^{(t)} + \Delta\alpha_{[k]}) \\ &= \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \left(-\frac{\lambda}{2} \|A(\alpha^{(t)} + \Delta\alpha_{[k]})\|^2 - \frac{1}{n} \sum_{i=1}^n \ell_i^*(-(\alpha^{(t)} + \Delta\alpha_{[k]})_i) \right)\end{aligned}$$

- we cannot solve the subproblem as it depends on $\alpha^{(t)}$ and A
- if we know $w^{(t)} = A\alpha^{(t)}$ then

$$(\Delta\alpha^*)_{[k]}^{(t)} = \arg \max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \left(-\frac{\lambda}{2} \|w^{(t)} + A\Delta\alpha_{[k]}\|^2 - \frac{1}{n} \sum_{i \in \mathcal{P}_k} \ell_i^*(-(\alpha^{(t)} + \Delta\alpha_{[k]})_i) \right)$$

- if we know $w^{(t)}$ we can compute $(\Delta\alpha^*)_{[k]}^{(t)}$

The CoCoA Framework

Communication-Efficient Distributed Dual Coordinate Ascent

Input: $T \geq 1$

Data: $\{(x_i, y_i)\}_{i=1}^n$ distributed over K machines

Initialize: $\alpha_{[k]}^{(0)} \leftarrow 0$ for all machines k , and $w^{(0)} \leftarrow 0$

for $t = 1, 2, \dots, T$

for all machines $k = 1, 2, \dots, K$ **in parallel**

Solve local problem approximately to obtain $\Delta\alpha_{[k]}$

computation

$$\alpha_{[k]}^{(t)} \leftarrow \alpha_{[k]}^{(t-1)} + \frac{1}{K} \Delta\alpha_{[k]}$$

$$\Delta w_k \leftarrow \frac{1}{K} A \Delta\alpha_{[k]}$$

$$\text{reduce } w^{(t)} \leftarrow w^{(t-1)} + \sum_{k=1}^K \Delta w_k$$

communication

- The performance of this methods (in worst case) can be the same as if we **randomly pick** k and solve corresponding subproblem and replace $\frac{1}{K}$ by 1
- How **accurately** do we need to solve the local sub-problem?
- How to change the local problem to avoid averaging (e.g. just to add local solutions)?
- Can we prove it will be better?

Smarter Subproblem

Local Subproblem for CoCoA⁺

$$\max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}; w^{(t)}) \quad (2)$$

where

$$\begin{aligned} \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}; w^{(t)}) = & -\frac{1}{n} \sum_{i \in \mathcal{P}_k} \ell_i^*(-(\alpha_{[k]}^{(t)} + \Delta\alpha_{[k]})_i) \\ & - \frac{1}{K} \frac{\lambda}{2} \|w^{(t)}\|^2 - \lambda (w^{(t)})^T A \Delta\alpha_{[k]} \\ & - \frac{\lambda}{2} \sigma' \|A \Delta\alpha_{[k]}\|^2. \end{aligned}$$

Compare with:

$$\max_{\Delta\alpha_{[k]} \in \mathbb{R}^n} \left(-\frac{\lambda}{2} \|w^{(t)} + A \Delta\alpha_{[k]}\|^2 - \frac{1}{n} \sum_{i \in \mathcal{P}_k} \ell_i^*(-(\alpha^{(t)} + \Delta\alpha_{[k]})_i) \right) \quad (3)$$

If $\sigma' = 1$ then the optimal solutions of (2) and (3) coincide.

The CoCoA⁺ Framework

Communication-Efficient Distributed Dual Coordinate Ascent

Input: $T \geq 1$, $\gamma \in [\frac{1}{K}, 1]$, $\sigma' \in [1, \infty)$

Data: $\{(x_i, y_i)\}_{i=1}^n$ distributed over K machines

Initialize: $\alpha_{[k]}^{(0)} \leftarrow 0$ for all machines k , and $w^{(0)} \leftarrow 0$

for $t = 1, 2, \dots, T$

for all machines $k = 1, 2, \dots, K$ **in parallel**

 approximately $\max \mathcal{G}^{\sigma'}(\Delta\alpha_{[k]}; w^{(t)})$ to obtain $\Delta\alpha_{[k]}$ computation

$$\alpha_{[k]}^{(t)} \leftarrow \alpha_{[k]}^{(t-1)} + \gamma \Delta\alpha_{[k]}$$

$$\Delta w_k \leftarrow \gamma A \Delta\alpha_{[k]}$$

 reduce $w^{(t)} \leftarrow w^{(t-1)} + \sum_{k=1}^K \Delta w_k$ communication

- If $\gamma = \frac{1}{K}$ we obtain CoCoA
- If $\gamma = \frac{1}{K}$ then $\sigma' = 1$ is "safe" value
- What about another values of γ ? (we want $\gamma = 1$)

CoCoA⁺ Parameters - σ' and γ

- σ' measures the difficulty of the given data partition
- it must be chosen not smaller than

$$\sigma' \geq \sigma'_{\min} \stackrel{\text{def}}{=} \gamma \max_{\alpha \in \mathbf{R}^n} \frac{\|A\alpha\|^2}{\sum_{k=1}^K \|A\alpha_{[k]}\|^2} \quad (4)$$

Lemma

For any $\alpha \in \mathbf{R}^n$ ($\alpha \neq \mathbf{0}$) we have

$$\frac{\|A\alpha\|^2}{\sum_{k=1}^K \|A\alpha_{[k]}\|^2} \leq K$$

- We can take the safe value $\sigma' = K \cdot \gamma$
Again: if $\gamma = \frac{1}{K}$ then $\sigma' = K \cdot \frac{1}{K} = 1$ is a safe value
- New: if $\gamma = 1$ then $\sigma' = K \cdot 1 = K$ is a safe value

How Accurately?

Assumption: Θ -approximate solution

We assume that there exists $\Theta \in [0, 1]$ such that $\forall k \in [K]$, the local solver at any iteration t produces a **(possibly) randomized** approximate solution $\Delta\alpha_{[k]}$, which satisfies

$$\mathbf{E}[\mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}^*, w) - \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}, w)] \leq \Theta \left(\mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}^*, w) - \mathcal{G}_k^{\sigma'}(\mathbf{0}, w) \right), \quad (5)$$

where

$$\Delta\alpha^* \in \arg \min_{\Delta\alpha \in \mathbf{R}^n} \sum_{k=1}^K \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}, w). \quad (6)$$

- because the subproblem is **not really** what one wants to solve, therefore in practise $\Theta \approx 0.9$ (depending on the cluster and problem)
- what about convergence guarantees?
- how to get Θ approximate solution?

Iteration Complexity - Smooth Loss

Theorem

Assume the loss functions ℓ_i are $(1/\mu)$ -smooth, for $i \in \{1, 2, \dots, n\}$. We define

$$\sigma_k \stackrel{\text{def}}{=} \max_{\alpha_{[k]} \in \mathbf{R}^n} \frac{\|A\alpha_{[k]}\|^2}{\|\alpha_{[k]}\|^2} \leq |\mathcal{P}_k| \quad (7)$$

and $\sigma_{\max} = \max_{k \in [K]} \sigma_k$.

Then after T iterations of CoCoA⁺, with

$$T \geq \frac{1}{\gamma(1-\Theta)} \frac{\lambda\mu n + \sigma_{\max}\sigma'}{\lambda\mu n} \log \frac{1}{\epsilon},$$

it holds that $\mathbf{E}[D(\alpha^*) - D(\alpha^T)] \leq \epsilon$.

Furthermore, after T iterations with

$$T \geq \frac{1}{\gamma(1-\Theta)} \frac{\lambda\mu n + \sigma_{\max}\sigma'}{\lambda\mu n} \log \left(\frac{1}{\gamma(1-\Theta)} \frac{\lambda\mu n + \sigma_{\max}\sigma'}{\lambda\mu n} \frac{1}{\epsilon} \right),$$

we have the expected duality gap

$$\mathbf{E}[P(w(\alpha^{(T)})) - D(\alpha^{(T)})] \leq \epsilon.$$

Averaging vs. Adding

The leading term is $\frac{1}{\gamma(1-\Theta)} \frac{\lambda\mu n + \sigma_{\max}\sigma'}{\lambda\mu n}$. Let us assume that $\forall k : |\mathcal{P}_k| = \frac{n}{K}$

Averaging

$$\gamma = \frac{1}{K}$$
$$\sigma' = 1$$

$$\frac{\textcolor{red}{K}}{1-\Theta} \frac{\lambda\mu n + \frac{n}{K}}{\lambda\mu n}$$

$$\frac{1}{1-\Theta} \frac{\lambda\mu \textcolor{red}{K} + 1}{\lambda\mu}$$

Adding

$$\gamma = 1$$
$$\sigma' = K$$

$$\frac{1}{1-\Theta} \frac{\lambda\mu n + \frac{n}{K} \textcolor{red}{K}}{\lambda\mu n}$$

$$\frac{1}{1-\Theta} \frac{\lambda\mu + 1}{\lambda\mu}$$

Note: this is in the worst case (for the worst case example)

Iteration Complexity - General Convex Loss

Theorem

Consider CoCoA⁺ starting with $\alpha^0 = \mathbf{0} \in \mathbf{R}^n$ and $\forall i \in \{1, 2, \dots, n\} : \ell_i(\cdot)$ be L -Lipschitz continuous and $\epsilon > 0$ be the desired duality gap. Then after T iterations, where

$$\begin{aligned} T &\geq T_0 + \max\left\{\left\lceil \frac{1}{\gamma(1-\Theta)} \right\rceil, \frac{4L^2\sigma\sigma'}{\lambda n^2\epsilon\gamma(1-\Theta)}\right\}, \\ T_0 &\geq t_0 + \left(\frac{2}{\gamma(1-\Theta)} \left(\frac{8L^2\sigma\sigma'}{\lambda n^2\epsilon} - 1\right)\right)_+, \\ t_0 &\geq \max(0, \left\lceil \frac{1}{\gamma(1-\Theta)} \log\left(\frac{2\lambda n^2(D(\alpha^*) - D(\alpha^0))}{4L^2\sigma\sigma'}\right) \right\rceil), \end{aligned}$$

we have that the expected duality gap satisfies $\mathbf{E}[P(w(\bar{\alpha})) - D(\bar{\alpha})] \leq \epsilon$, at the averaged iterate

$$\bar{\alpha} := \frac{1}{T-T_0} \sum_{t=T_0+1}^{T-1} \alpha^{(t)},$$

where $\sigma = \sum_{k=1}^K |\mathcal{P}_k| \sigma_k$.

SDCA as a Local Solver

SDCA

- 1: **Input:** $\alpha_{[k]}, w = w(\alpha)$
- 2: **Data:** Local $\{(x_i, y_i)\}_{i \in \mathcal{P}_k}$
- 3: **Initialize:** $\Delta\alpha_{[k]}^0 = 0 \in \mathbb{R}^n$
- 4: **for** $h = 0, 1, \dots, H-1$ **do**
- 5: choose $i \in \mathcal{P}_k$ uniformly at random
- 6: $\delta_i^* = \arg \max_{\delta_i \in \mathbb{R}} \mathcal{G}_k^{\sigma'}(\Delta\alpha_{[k]}^h + \delta_i e_i, w)$
- 7: $\Delta\alpha_{[k]}^{(h+1)} = \Delta\alpha_{[k]}^{(h)} + \delta_i^* e_i$
- 8: **end for**
- 9: **Output:** $\Delta\alpha_{[k]}^{(H)}$

Theorem

Assume the functions ℓ_i are $(1/\mu)$ -smooth for $i \in \{1, 2, \dots, n\}$. If

$$H \geq n_k \frac{\sigma' + \lambda n \mu}{\lambda n \mu} \log \frac{1}{\Theta} \quad (8)$$

then SDCA will produce a Θ -approximate solution.

Total Runtime

- To get ϵ accuracy we need

$$\mathcal{O}\left(\frac{1}{1 - \Theta} \log \frac{1}{\epsilon}\right)$$

- Recall $\Theta = \left(1 - \frac{\lambda n \gamma}{1 + \lambda n \gamma} \frac{K}{n}\right)^H$





Let

- τ_o be the duration of communication per iteration
- τ_c be the duration of **ONE** coordinate update during the inner iteration

Total runtime

$$\mathcal{O}\left(\frac{1}{1 - \Theta} (\tau_o + H \tau_c)\right) = \mathcal{O}\left(\frac{1}{1 - \Theta} \left(1 + H \underbrace{\frac{\tau_c}{\tau_o}}_{r_{c/o}}\right)\right)$$

$$H(\tau_c/\tau_o), \Theta(\tau_c/\tau_o)$$



Apache Spark


- is a fast and general engine for large-scale data processing
- runs on Hadoop, Mesos, standalone, or in the cloud. It can access diverse data sources including HDFS, Cassandra, HBase, S3
- **is slower than our C++ code** for CoCoA⁺
- we run it on Amazon Elastic Compute Cloud (Amazon EC2)

Numerical Experiments

Datasets

Dataset	Training (n)	Features (d)	$nnz/(n \cdot d)$	λ	Workers (K)
cov	522,911	54	22.22%	$1e-6$	4
rcv1	677,399	47,236	0.16%	$1e-6$	8
imagenet	32,751	160,000	100%	$1e-5$	32

Dependence of Primal Suboptimality on H



Comparison with Different Algorithms

- COCOA

Comparison with Different Algorithms

- COCOA

- minibatch-CD

Tianbao Yang. *Trading Computation for Communication: Distributed Stochastic Dual Coordinate Ascent*. In NIPS 2013.

Martin Takáč, Avleen Bijral, Peter Richtárik, and Nathan Srebro. *Mini-Batch Primal and Dual Methods for SVMs*. In ICML, March 2013.

Comparison with Different Algorithms

- COCOA
- minibatch-CD
Tianbao Yang. *Trading Computation for Communication: Distributed Stochastic Dual Coordinate Ascent*. In NIPS 2013.
Martin Takáč, Avleen Bijral, Peter Richtárik, and Nathan Srebro. *Mini-Batch Primal and Dual Methods for SVMs*. In ICML, March 2013.
- local-SGD
Shai Shalev-Shwartz, Yoram Singer, Nathan Srebro, and Andrew Cotter. *Pegasos: Primal Estimated Sub-Gradient Solver for SVM*. Mathematical Programming, 127(1):3–30, October 2010.

Comparison with Different Algorithms

- COCOA

- minibatch-CD

Tianbao Yang. *Trading Computation for Communication: Distributed Stochastic Dual Coordinate Ascent*. In NIPS 2013.




Martin Takáč, Avleen Bijral, Peter Richtárik, and Nathan Srebro. *Mini-Batch Primal and Dual Methods for SVMs*. In ICML, March 2013.




- local-SGD

Shai Shalev-Shwartz, Yoram Singer, Nathan Srebro, and Andrew Cotter. *Pegasos: Primal Estimated Sub-Gradient Solver for SVM*. Mathematical Programming, 127(1):3–30, October 2010.





- batch-SGD

Shai Shalev-Shwartz, Yoram Singer, Nathan Srebro, and Andrew Cotter. *Pegasos: Primal Estimated Sub-Gradient Solver for SVM*. Mathematical Programming, 127(1):3–30, October 2010.












CoCoA vs. CoCoA⁺





CoCoA vs. CoCoA⁺



Scaling up



Effect of σ'



References

- ① Chenxin Ma, Virginia Smith, Martin Jaggi, Michael I. Jordan, Peter Richtárik and Martin Takáč: *Adding vs. Averaging in Distributed Primal-Dual Optimization*, ICML 2015.
- ② Martin Jaggi, Virginia Smith, Martin Takáč, Jonathan Terhorst, Thomas Hofmann and Michael I. Jordan: *Communication-Efficient Distributed Dual Coordinate Ascent*, NIPS 2014.
- ③ Richtárik, P. and Takáč, M.: *On optimal probabilities in stochastic coordinate descent methods*, arXiv:1310.3438, 2013.
- ④ Richtárik, P. and Takáč, M.: *Parallel coordinate descent methods for big data optimization*, arXiv:1212.0873, 2012.
- ⑤ Richtárik, P. and Takáč, M.: *Iteration complexity of randomized block-coordinate descent methods for minimizing a composite function*, Mathematical Programming, 2012.
- ⑥ Takáč, M., Bijral, A., Richtárik, P. and Srebro, N.: *Mini-batch primal and dual methods for SVMs*, In ICML, 2013.
- ⑦ Qu, Z., Richtárik, P. and Zhang, T.: *Randomized dual coordinate ascent with arbitrary sampling*, arXiv:1411.5873, 2014.
- ⑧ Qu, Z., Richtárik, P., Takáč, M. and Fercoq, O.: *SDNA: Stochastic Dual Newton Ascent for Empirical Risk Minimization*, arXiv:1502.02268, 2015.
- ⑨ Tappenden, R., Takáč, M. and Richtárik, P., *On the Complexity of Parallel Coordinate Descent*, arXiv: 1503.03033, 2015.